


yaml goes in, pods come out.

declarative


description of an application



k8s cluster

"they go do stuff"

"The Brains"



control plane

- scheduling
- monitoring

usually there are 3 masters.

- you should have an odd # of masters

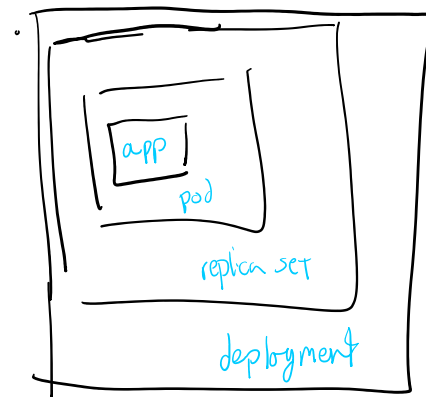
◦ only one master is a "leader" at a time

- the others are followers
- followers proxy reqs to the leader

... 11 ... masters

"reconciliation loop"

- control plane constantly checking desired state vs observed state



levels of k8s units of work

- users interface with the masters via k8s' api server

- spec once, deploy many

The API Server

- everything you interact w/ is an api object
- we feed yaml or json to the api to do stuff
- lives on the control plane
- kubectl talks to this when issuing commands